

The Llama fine-tuning libraries

FINE-TUNING WITH LLAMA 3




Francesca Donadoni

Curriculum Manager, DataCamp

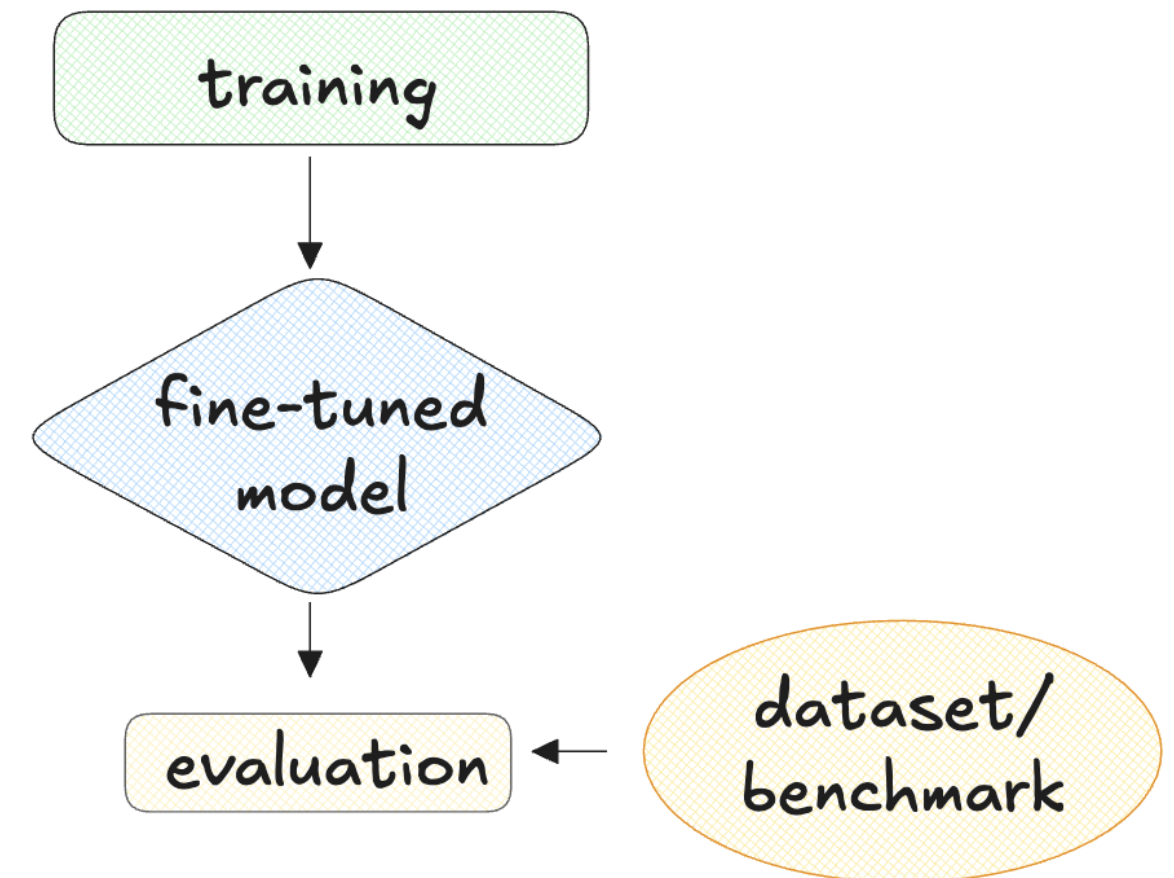
When to use fine-tuning

- Pre-trained model
- Uses specialized data
- Improve accuracy
- Reduce bias
- Improve knowledge base



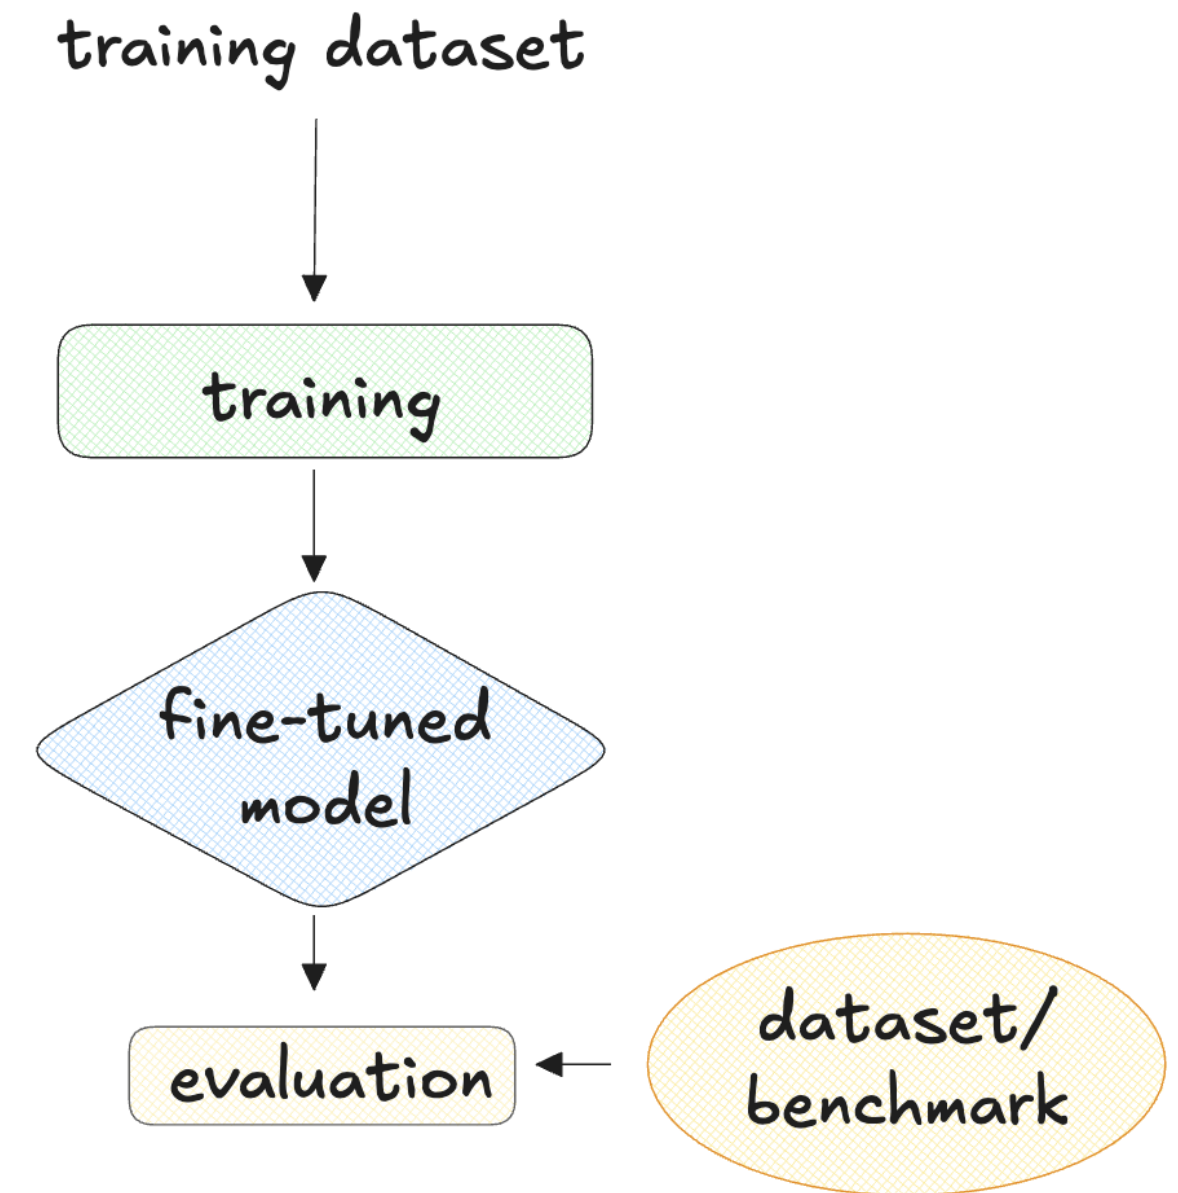
How to use fine-tuning

- Quality of the data
- Model's capacity
- Task definition
- Fine-tuning process



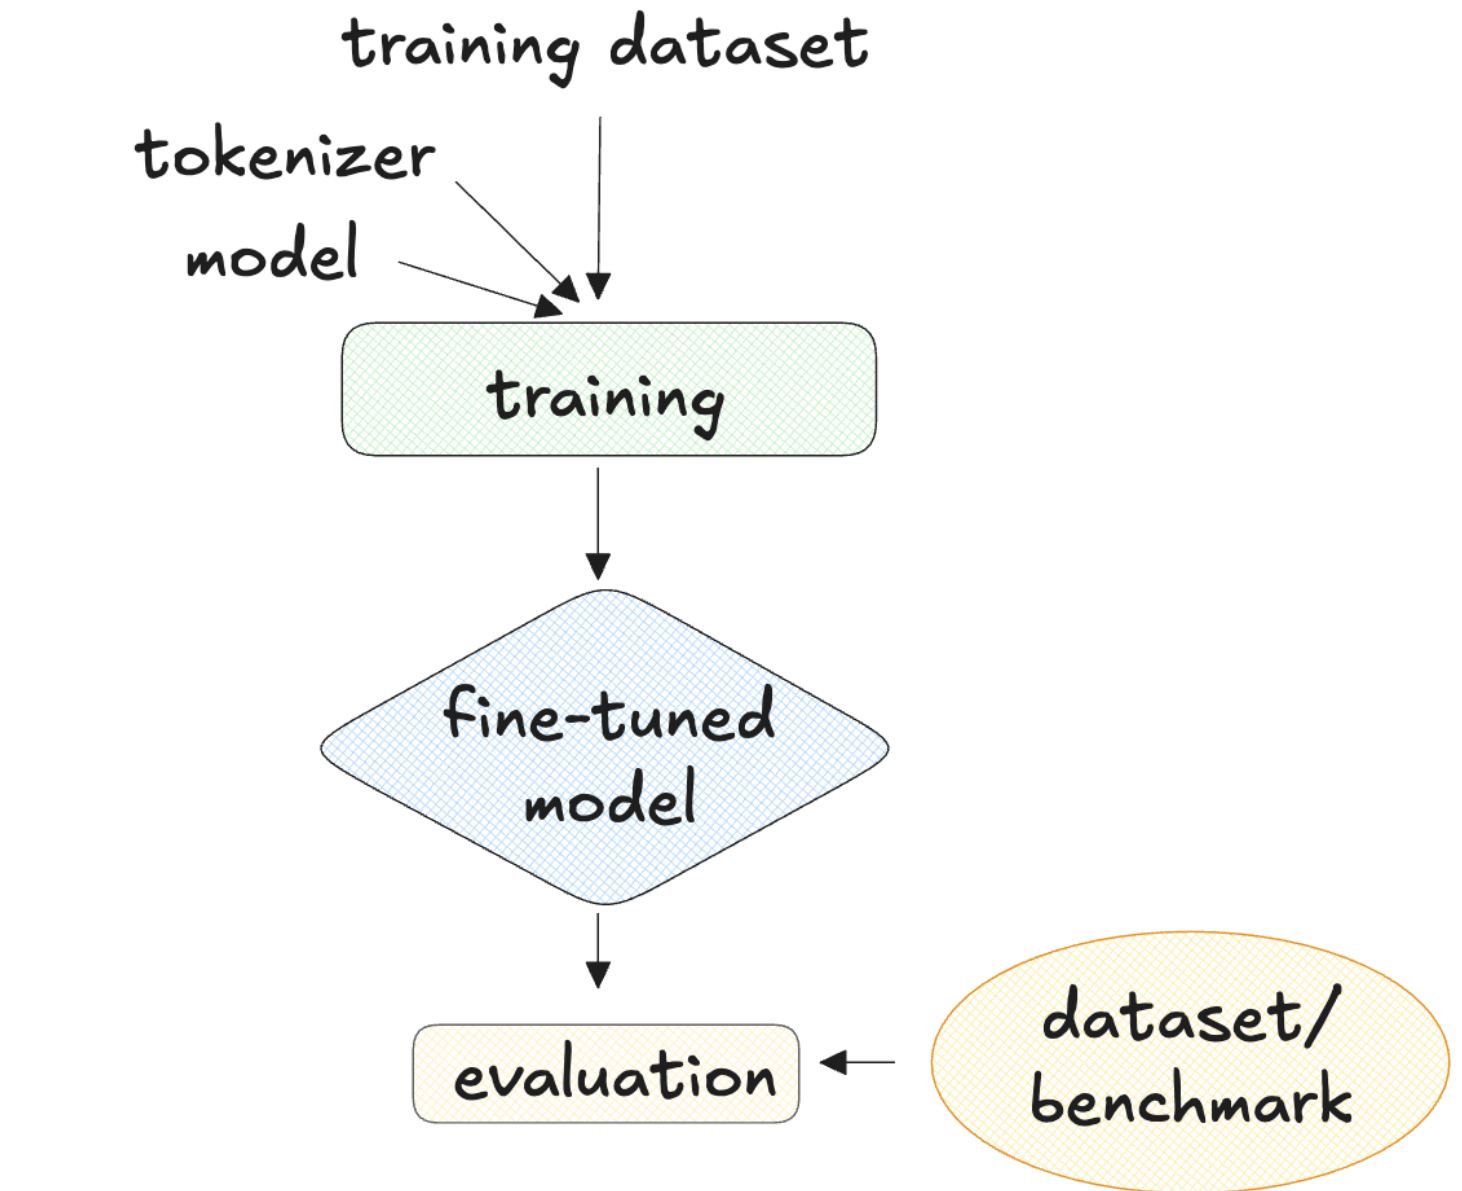
How to use fine-tuning

- Quality of the data
- Model's capacity
- Task definition
- Fine-tuning process



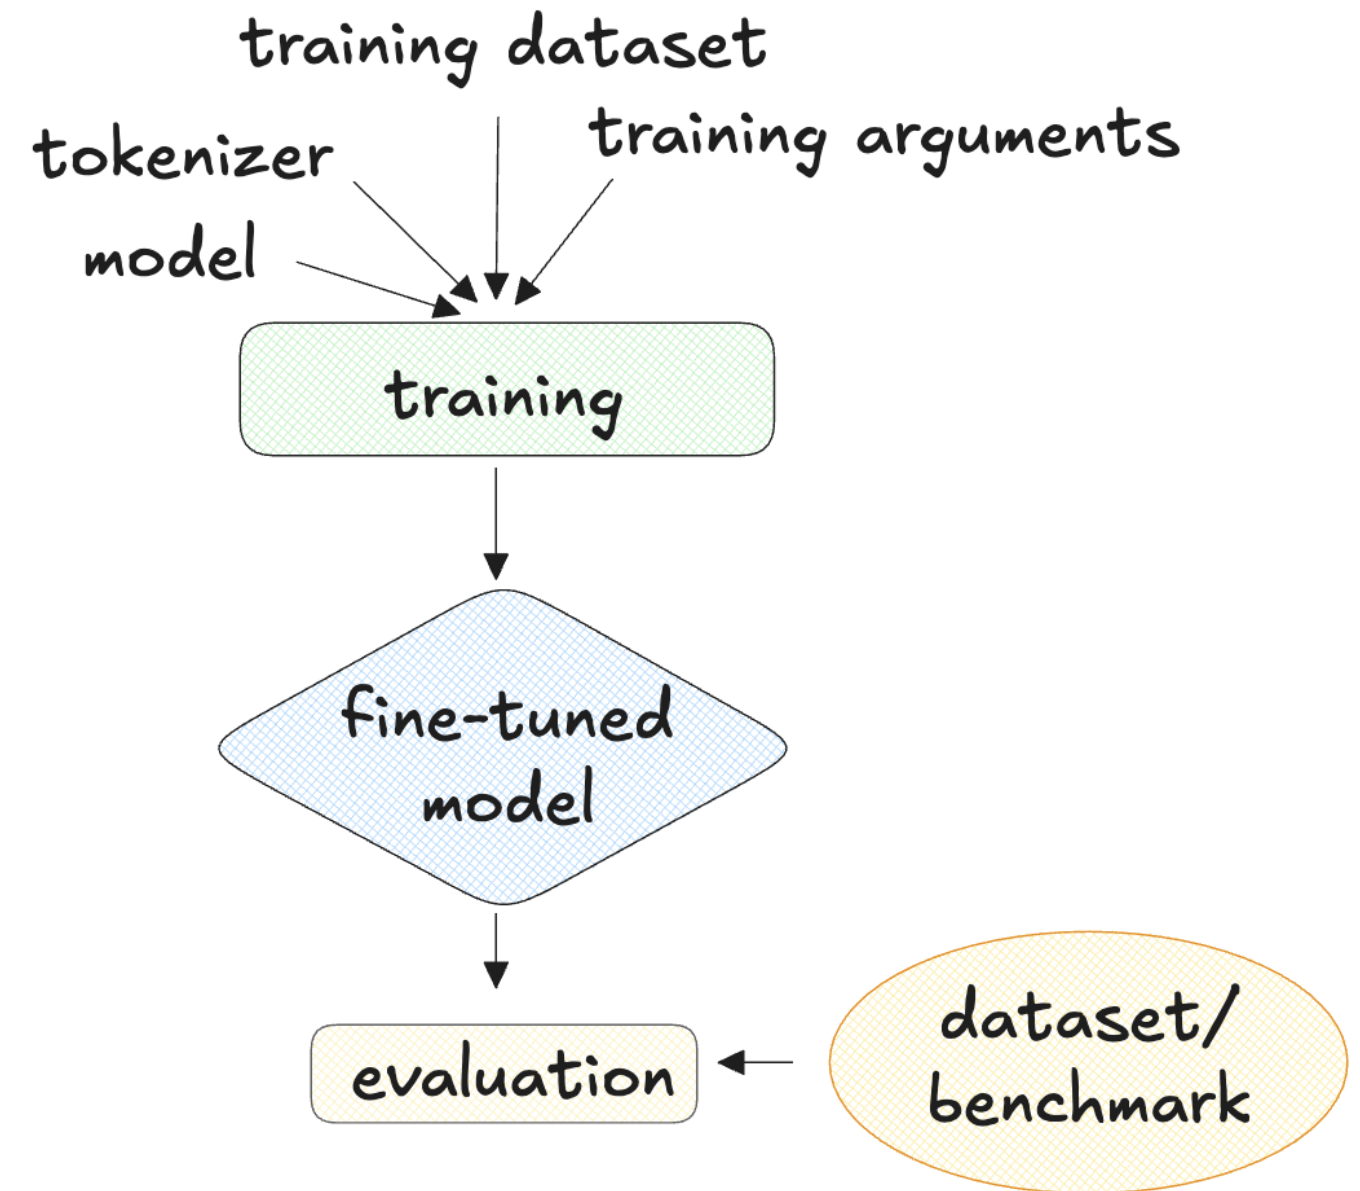
How to use fine-tuning

- Quality of the data
- Model's capacity
- Task definition
- Fine-tuning process



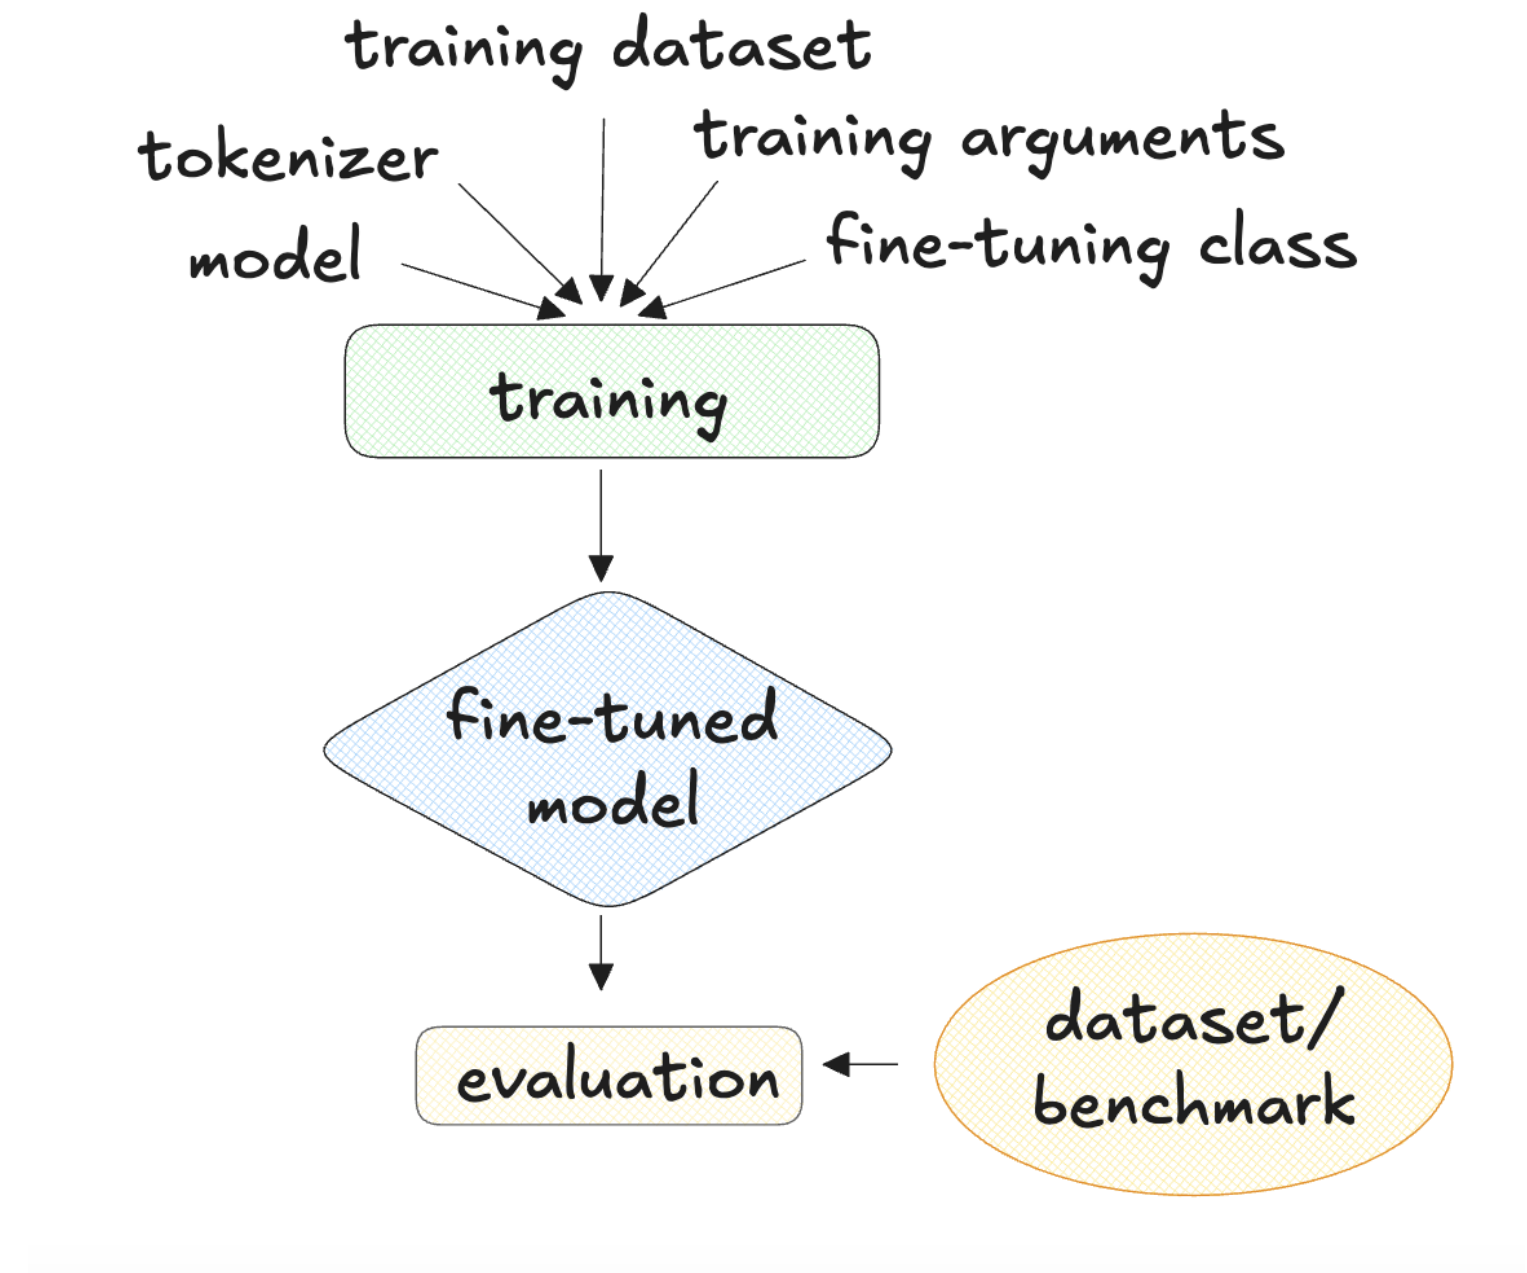
How to use fine-tuning

- Quality of the data
- Model's capacity
- Task definition
- Fine-tuning process



How to use fine-tuning

- Quality of the data
- Model's capacity
- Task definition
- Fine-tuning process
- New model
- Evaluation



The Llama fine-tuning libraries

- ☐ Several libraries for fine-tuning
- ☐ TorchTune for Llama fine-tuning
- ☐ Launching a fine-tuning task with TorchTune

Options for Llama fine-tuning


- **TorchTune**

- Based on configurable templates
- Ideal for: **scaling quickly**




- **Unsloth**

- Efficient memory usage
- Ideal for: **limited hardware**



- **SFTTrainer from Hugging Face**

- Access to other LLMs
- Ideal for: **fine-tuning multiple models**



- **Axolotl**

- Modular approach
- Ideal for: **no extensive reconfiguration**



TorchTune and the recipes for fine-tuning

- TorchTune recipes:
 - Modular templates
 - Configurable to be adapted to different projects
 - Keep code organized
 - Ensure reproducibility



TorchTune list

- Run from a terminal
- Environment with Python
- Install TorchTune

```
pip3 install torchtune
```

- List available recipes

```
tune ls
```

- ! if using IPython

```
!tune ls
```

TorchTune list

```
!tune ls
```

- Output:

RECIPE	CONFIG
full_finetune_single_device	llama3/8B_full_single_device
	llama3_1/8B_full_single_device
	llama3_2/1B_full_single_device
	llama3_2/3B_full_single_device
full_finetune_distributed	llama3/8B_full
	llama3_1/8B_full
	llama3_2/1B_full
	...

TorchTune run

- Use recipe + `--config` + configuration
- Run fine-tuning

```
tune run full_finetune_single_device --config \
llama3_1/8B_lora_single_device
```

- Parameters `device=cpu` or `device=cuda`
- `epochs=<int>` (`<int>` is 0 or a positive integer)

Let's practice!

FINE-TUNING WITH LLAMA 3

Preprocessing data for fine-tuning

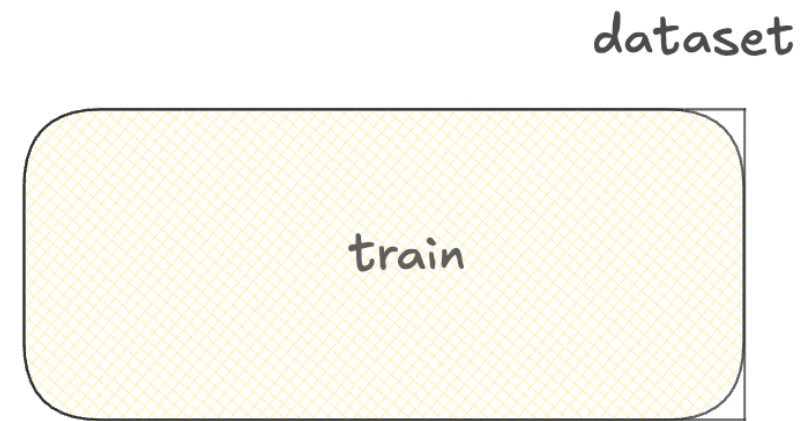
FINE-TUNING WITH LLAMA 3



Francesca Donadoni
Curriculum Manager, DataCamp

Using datasets for fine-tuning

- Quality of the data is key
- Training Set:
 - For model training
 - Majority of the data




Using datasets for fine-tuning

- Quality of the data is key
- **Training Set:**
 - For model training
 - Majority of the data
- **Validation Set:**
 - For selecting the best model version



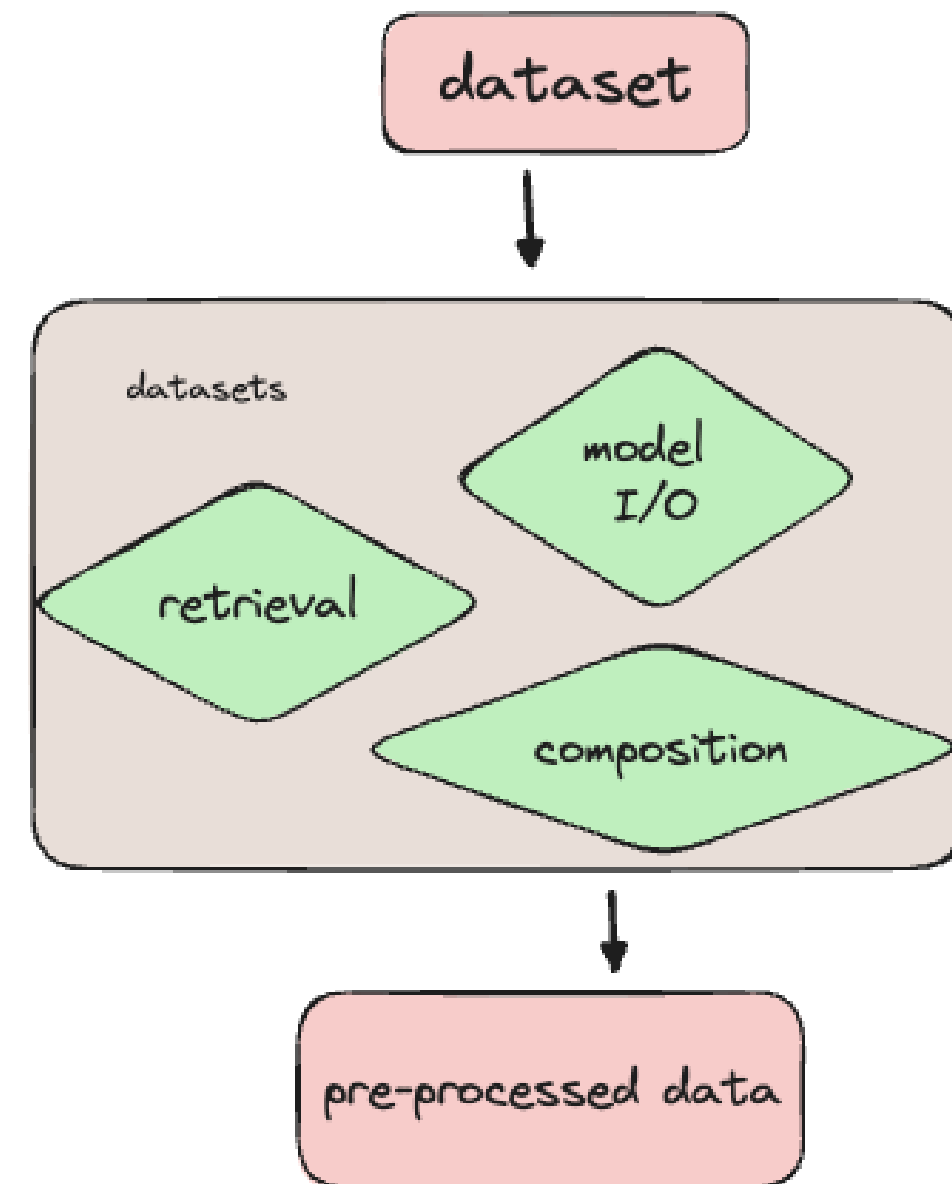
Using datasets for fine-tuning

- Quality of the data is key
- **Training Set:**
 - For model training
 - Majority of the data
- **Validation Set:**
 - For selecting the best model version
- **Test Set:**
 - For evaluating model's performance



Preparing data using the datasets library

- Datasets library
- Preprocessing
- Split
- Load
- Manage memory



Loading a customer service dataset

```
from datasets import load_dataset
ds = load_dataset(
    'bitext/Bitext-customer-support-llm-chatbot-training-dataset',
    split="train"
)
print(ds.column_names)
```

```
['flags', 'instruction', 'category', 'intent', 'response']
```

Peeking into the data

```
import pprint
pprint.pprint(ds[0])
```

```
{'category': 'ORDER',
 'flags': 'B',
 'instruction': 'question about cancelling order {{Order Number}}',
 'intent': 'cancel_order',
 'response': "I've understood you have a question regarding canceling order "
              "{{Order Number}}, and I'm here to provide you with the "
              'information you need. Please go ahead and ask your question, and '
              "I'll do my best to assist you."}
```

Filtering the dataset

```
from datasets import load_dataset, Dataset

ds = load_dataset(
    'bitext/Bitext-customer-support-llm-chatbot-training-dataset',
    split="train")
print(ds.shape)
```

```
(26872, 5)
```

```
first_thousand_points = ds[:1000]
ds = Dataset.from_dict(first_thousand_points)
```


Preprocessing the dataset

```
def merge_example(row):  
    row['conversation'] = f"Query: {row['instruction']}\nResponse: {row['response']}"  
    return row  
ds = ds.map(merge_example)  
print(ds[0]['conversation'])
```

```
Query: question about cancelling order {{Order Number}}  
Response: I've understood you have a question regarding canceling order {{Order Number}},  
and I'm here to provide you with the information you need. Please go ahead and ask your  
question, and I'll do my best to assist you.
```

Saving the preprocessed dataset

```
ds.save_to_disk("preprocessed_dataset")
```

```
Saving the dataset (1/1 shards): 100%  
26872/26872 [00:00<00:00, 383823.33 examples/s]
```

```
from datasets import load_from_disk  
ds_preprocessed = load_from_disk("preprocessed_dataset")
```

Using Hugging Face datasets with TorchTune

- Can use Hugging Face dataset with TorchTune
- Set a dataset path and configurations

```
tune run full_finetune_single_device --config llama3/8B_full_single_device \
dataset=preprocessed_dataset dataset.split=train
```

Let's practice!

FINE-TUNING WITH LLAMA 3

Fine-tuning with TorchTune



FINE-TUNING WITH LLAMA 3



Francesca Donadoni
Curriculum Manager, DataCamp

The components of TorchTune fine-tuning

- **Model**
 - Defines the architecture and pre-trained weights to fine-tune
 - Different versions and number of parameters available
- **Dataset**
 - Specifies the data used for training
- **Recipe**
 - Central configuration file combining the model, dataset, and training parameters
 - Ensures consistency and reproducibility



The components of TorchTune fine-tuning

- **Model**

- `!tune ls`



```
llama3/8B_full  
llama3_1/8B_full  
llama3_2/1B_full ...
```

- **Dataset**

- `ds.save_to_disk("new_dataset")`

- **Recipe**

- `custom_recipe.yaml`



The components of a TorchTune recipe

- General Settings and output directory
 - Batch size, device and epochs
- Model
 - Specifies architecture and model configurations
- Optimizer
 - Includes learning rate
- Dataset
 - Defines preprocessing and dataset path

```
batch_size: 4
device: cuda
epochs: 20
output_dir: /tmp/full-llama3.2-finetune

model:
  _component_:
    torchtune.models.llama3_2.llama3_2_1b

optimizer:
  _component_: bitsandbytes.optim.PagedAdamW8bit
  lr: 2.0e-05

dataset:
  _component_: torchtune.datasets.alpaca_dataset
```

Configuring TorchTune recipes

- More parameters available
- Configurable in Python using `yaml`

```
import yaml
config_dict = {"batch_size": 4,
               "device": "cuda",
               "model": {
                   "_component_": "torchtune.models.llama3_2.llama3_2_1b"
               },
               ...
}
yaml_file_path = "custom_recipe.yaml"
with open(yaml_file_path, "w") as yaml_file:
    yaml.dump(config_dict, yaml_file)
```

Running custom fine-tuning

```
tune run --config custom_recipe.yaml
```

```
INFO:torchutils.logging:Running  
Writing logs to /tmp/full-llama3.2-finetune/log_1732815689.txt  
INFO:torchutils.logging:Model is initialized with precision torch.bfloat16.  
INFO:torchutils.logging:Tokenizer is initialized from file.  
1|52|Loss: 2.3697006702423096: 0%|? | 52/25880
```

- Saved logs
- Successful initialization
- Epoch and step count progress
- Loss metrics

Let's practice!

FINE-TUNING WITH LLAMA 3