


API versioning and documentation

DEPLOYING AI INTO PRODUCTION WITH FASTAPI



Matt Eckerle

Software and Data Engineering Leader

Why API versioning?

- API endpoints as menu items
- Keep old customers happy
- Some customers want new options
- **Iterate without impacting existing customers**



API for cloud AI jobs

```
from pydantic import BaseModel

class AIJobV1(BaseModel):
    job_name: str
    data: bytes
```

```
class AIJobV1(BaseModel):
    job_name: str
    data: bytes
    config: dict
```

 FastAPI



Versioned endpoints

```
from pydantic import BaseModel

class AIJobV1(BaseModel):
    job_name: str
    data: bytes

class AIJobV2(BaseModel):
    job_name: str
    data: bytes
    config: dict
```

```
from fastapi import FastAPI

app = FastAPI()

@app.post("/v1/ai-job")
def ai_job_v1(job: AIJobV1):
    ...

@app.post("/v2/ai-job")
def ai_job_v2(job: AIJobV2):
    ...
```

Reasons to update endpoint version

- Breaking change in schema
- Change in underlying function
 - Updated model code
 - Updated model training set
 - Updated pre/post processing

Iteration with optional fields

- Versioning is not always required to iterate
- Optional fields can support additional data without breaking schemas

```
from pydantic import BaseModel
```

```
class AIJobV1(BaseModel):
```

```
    job_name: str
```

```
    data: bytes
```

```
from typing import Optional
```

```
class AIJobV1(BaseModel):
```

```
    job_name: str
```

```
    data: bytes
```

```
    config: Optional[dict]
```

Documenting APIs with Swagger

- Standard tool for API documentation
 - Keeps track of endpoints and versions
 - Built on OpenAPI standard metadata




Swagger™



OPENAPI

Swagger UI



FastAPI 0.1.0 OAS 3.1

/openapi.json

AI Job API

default

POST	/v1/ai-job	Ai Job V1	⌵
POST	/v2/ai-job	Ai Job V2	⌵

Schemas

AIJobV1	>	Expand all	object
AIJobV2	>	Expand all	object
HTTPValidationError	>	Expand all	object
ValidationError	>	Expand all	object

Swagger UI for an endpoint

POST

/v1/ai-job

Ai Job V1

^

Parameters

Try it out

No parameters

Request body

required

application/json

▼

Example Value

Schema

```
{
  "job_name": "string",
  "data": "string"
}
```

Responses

Code	Description	Links
200	Successful Response	No links

Media type

application/json

▼

Controls Accept header.

Using FastAPI's description field

```
from fastapi import FastAPI

app = FastAPI(
    description="AI Job API"
)
```

FastAPI 0.1.0 OAS 3.1

/openapi.json


AI Job API

Let's practice!

DEPLOYING AI INTO PRODUCTION WITH FASTAPI

Advanced input validation and error handling

DEPLOYING AI INTO PRODUCTION WITH FASTAPI



Matt Eckerle

Software and Data Engineering Leader

Why we need advanced input

- API for restaurant orders
- Variable number of items

```
class Order(BaseModel):  
    item1: str  
    item2: str  
    item3: str
```



Nested Pydantic models

```
from pydantic import BaseModel

class Foo(BaseModel):
    count: int

class Bar(BaseModel):
    foo: Foo
```

```
>>> m = Bar(foo={'count': 4})
>>> print(m)
foo=Foo(count=4)
```

```
from pydantic import BaseModel
from typing import List

class OrderItem(BaseModel):
    name: str
    quantity: int

class RestaurantOrder(BaseModel):
    customer_name: str
    items: List[OrderItem]
```

Custom model validators

```
from fastapi import FastAPI
from fastapi.exceptions import (
    RequestValidationError
)
from pydantic import (
    BaseModel,
    model_validator,
)
from typing import List

class OrderItem(BaseModel):
    name: str
    quantity: int
```

```
class RestaurantOrder(BaseModel):
    customer_name: str
    items: List[OrderItem]

    @model_validator(mode="after")
    def validate_after(self):
        if len(self.items) == 0:
            raise RequestValidationError(
                "No items in order!"
            )
        return self
```

```
{"detail": "No items in order!"}
```


Global exception handlers

```
from fastapi import FastAPI
from fastapi.exceptions import RequestValidationError
from fastapi.responses import PlainTextResponse

app = FastAPI()

@app.exception_handler(RequestValidationError)
async def validation_exception_handler(request, exc):
    msg = "Input validation error. See the documentation: http://127.0.0.1:8000/docs"
    return PlainTextResponse(msg, status_code=422)
```


```
Input validation error. See the documentation: http://127.0.0.1:8000/docs
```


Let's practice!

DEPLOYING AI INTO PRODUCTION WITH FASTAPI

Monitoring and logging

DEPLOYING AI INTO PRODUCTION WITH FASTAPI



Matt Eckerle

Software and Data Engineering Leader

Why monitoring and logging?

- Can't debug in production
- App supervisor needs a simple health check
- Logging key metrics over time



Setting up custom logging

- Load the uvicorn error logger
- Add custom logs to app startup
- Add custom logs to endpoints

```
from fastapi import FastAPI
import logging

logger = logging.getLogger(
    'uvicorn.error'
)

app = FastAPI()
logger.info("App is running!")

@app.get('/')
async def main():
    logger.debug('GET /')
    return 'ok'
```

Logging a when a model is loaded

```
from fastapi import FastAPI
import logging
import joblib

logger = logging.getLogger('uvicorn.error')

model = joblib.load('penguin_classifier.pkl')
logger.info("Penguin classifier loaded successfully.")

app = FastAPI()
```

Logging process time with middleware

```
from fastapi import FastAPI, Request
import logging
import time

logger = logging.getLogger('uvicorn.error')
app = FastAPI()

@app.middleware("http")
async def log_process_time(request: Request, call_next):
    start_time = time.perf_counter()
    response = await call_next(request)
    process_time = time.perf_counter() - start_time
    logger.info(f"Process time was {process_time} seconds.")
    return response
```

¹ <https://fastapi.tiangolo.com/tutorial/middleware/>

Setting the logging level

Log Level	Numeric Value
debug	10
info	20
warning	30
error	40
critical	50

```
uvicorn main:app --log-level debug
```


Monitoring

```
from fastapi import FastAPI

app = FastAPI()
@app.get("/health")
async def get_health():
    return {"status": "OK"}
```

- "I'm ok!"

Sharing model parameters with monitoring

```
from fastapi import FastAPI
import joblib

model = joblib.load(
    'penguin_classifier.pkl'
)

app = FastAPI()
@app.get("/health")
async def get_health():
    params = model.get_params()
    return {"status": "OK",
           "params": params}
```


- "I'm ok!"
- "Here are some fun facts about me!"

Let's practice!

DEPLOYING AI INTO PRODUCTION WITH FASTAPI

Wrap-up

DEPLOYING AI INTO PRODUCTION WITH FASTAPI



Matt Eckerle

Software and Data Engineering Leader

Introduction to FastAPI for Model Deployment

Chapter 1

- Basic GET and POST requests
- Loading a pre-trained model
- Running the `uvicorn` server
- Pydantic models for requests and responses



Integrating AI models

Chapter 2

- More structured input types
- Loading a pre-trained model in the app
- Structured prediction results



Field Validators

The `'Field'` function is used to customize and add metadata to fields of models



Custom Domain-Specific Validators

Create and apply custom validator functions




Validation Error Handling

Custom messages and user-friendly reporting

Securing and optimizing the API

Chapter 3

- API key authentication
- Rate limiting
- Async processing



API versioning, monitoring and logging

Chapter 4

- API versioning and documentation
- Advanced input validation and error handling
- Monitoring and logging



Congratulations!

DEPLOYING AI INTO PRODUCTION WITH FASTAPI