

DataFrame column operations

CLEANING DATA WITH PYSPARK



Mike Metzger
Data Engineering Consultant

DataFrame refresher

DataFrames:

- Made up of rows & columns
- Immutable
- Use various transformation operations to modify data

```
# Return rows where name starts with "M"
voter_df.filter(voter_df.name.like('M%'))
# Return name and position only
voters = voter_df.select('name', 'position')
```

Common DataFrame transformations

- Filter / Where

```
voter_df.filter(voter_df.date > '1/1/2019') # or voter_df.where(...)
```

- Select

```
voter_df.select(voter_df.name)
```

- withColumn

```
voter_df.withColumn('year', voter_df.date.year)
```

- drop

```
voter_df.drop('unused_column')
```

Filtering data

- Remove nulls
- Remove odd entries
- Split data from combined sources
- Negate with ~

```
voter_df.filter(voter_df['name'].isNotNull())  
voter_df.filter(voter_df.date.year > 1800)  
voter_df.where(voter_df['_c0'].contains('VOTE'))  
voter_df.where(~ voter_df._c1.isNotNull())
```

Column string transformations

- Contained in `pyspark.sql.functions`

```
import pyspark.sql.functions as F
```

- Applied per column as transformation

```
voter_df.withColumn('upper', F.upper('name'))
```

- Can create intermediary columns

```
voter_df.withColumn('splits', F.split('name', ' '))
```

- Can cast to other types

```
voter_df.withColumn('year', voter_df['_c4'].cast(IntegerType()))
```

ArrayType() column functions

Various utility functions / transformations to interact with ArrayType()

`.size(<column>)` - returns length of arrayType() column

`.getItem(<index>)` - used to retrieve a specific item at index of list column.

Let's practice!
CLEANING DATA WITH PYSPARK

Conditional DataFrame column operations

CLEANING DATA WITH PYSPARK



Mike Metzger
Data Engineering Consultant

Conditional clauses

Conditional Clauses are:

- Inline version of if / then / else
- `.when()`
- `.otherwise()`

Conditional example

```
.when(<if condition>, <then x>)
```

```
df.select(df.Name, df.Age, F.when(df.Age >= 18, "Adult"))
```

Name	Age	
Alice	14	
Bob	18	Adult
Candice	38	Adult

Another example

Multiple `.when()`

```
df.select(df.Name, df.Age,  
         .when(df.Age >= 18, "Adult")  
         .when(df.Age < 18, "Minor"))
```

Name	Age	
Alice	14	Minor
Bob	18	Adult
Candice	38	Adult

Otherwise

`.otherwise()` is like `else`

```
df.select(df.Name, df.Age,  
          .when(df.Age >= 18, "Adult")  
          .otherwise("Minor"))
```

Name	Age	
Alice	14	Minor
Bob	18	Adult
Candice	38	Adult

Let's practice!
CLEANING DATA WITH PYSPARK

User defined functions

CLEANING DATA WITH PYSPARK



Mike Metzger
Data Engineering Consultant

Defined...

User defined functions or UDFs

- Python method
- Wrapped via the `pyspark.sql.functions.udf` method
- Stored as a variable
- Called like a normal Spark function

Reverse string UDF

Define a Python method

```
def reverseString(mystr):  
    return mystr[::-1]
```

Wrap the function and store as a variable

```
udfReverseString = udf(reverseString, StringType())
```

Use with Spark

```
user_df = user_df.withColumn('ReverseName',  
                             udfReverseString(user_df.Name))
```


Argument-less example

```
def sortingCap():  
    return random.choice(['G', 'H', 'R', 'S'])  
udfSortingCap = udf(sortingCap, StringType())  
user_df = user_df.withColumn('Class', udfSortingCap())
```

Name	Age	Class
Alice	14	H
Bob	18	S
Candice	63	G

Let's practice!
CLEANING DATA WITH PYSPARK

Partitioning and lazy processing

CLEANING DATA WITH PYSPARK



Mike Metzger
Data Engineering Consultant

Partitioning

- DataFrames are broken up into partitions
- Partition size can vary
- Each partition is handled independently

Lazy processing

- Transformations are **lazy**
 - `.withColumn(...)`
 - `.select(...)`
- Nothing is actually done until an action is performed
 - `.count()`
 - `.write(...)`
- Transformations can be re-ordered for best performance
- Sometimes causes unexpected behavior

Adding IDs

Normal ID fields:

- Common in relational databases
- Most usually an integer increasing, sequential and unique
- Not very parallel

id	last name	first name	state
0	Smith	John	TX
1	Wilson	A.	IL
2	Adams	Wendy	OR

Monotonically increasing IDs

`pyspark.sql.functions.monotonically_increasing_id()`

- Integer (64-bit), increases in value, unique
- Not necessarily sequential (gaps exist)
- Completely parallel

id	last name	first name	state
0	Smith	John	TX
134520871	Wilson	A.	IL
675824594	Adams	Wendy	OR

Notes

Remember, Spark is *lazy*!

- Occasionally out of order
- If performing a join, ID may be assigned after the join
- Test your transformations

Let's practice!
CLEANING DATA WITH PYSPARK